

= ⚡ > ChatGPT 5.2 ✓

🔗 ⌂ ...

so the formula of perceptron trick

In ML literature, “perceptron trick” usually means the *kernelized perceptron* (analogous to the kernel trick in SVMs).

Below are the exact formulas, clean and mathematical.

1. Standard Perceptron (Primal Form)

Model

$$\hat{y} = \text{sign}(w^\top x + b)$$

Update rule (for a misclassified point (x_i, y_i) , $y_i \in \{-1, +1\}$)

$$\text{if } y_i(w^\top x_i + b) \leq 0 : \quad \begin{cases} w \leftarrow w + \eta y_i x_i \\ b \leftarrow b + \eta y_i \end{cases}$$

2. Dual Representation of Perceptron

Initialize $\alpha_i = 0 \ \forall i$

$$w = \sum_{i=1}^n \alpha_i y_i x_i$$

Decision function

$$f(x) = \sum_{i=1}^n \alpha_i y_i (x_i^\top x) + b$$

Update

$$\text{if } y_i f(x_i) \leq 0 : \quad \alpha_i \leftarrow \alpha_i + \eta$$

3. Perceptron Trick (Kernelized Perceptron)

Key substitution (the “trick”)

$$x_i^\top x \Rightarrow K(x_i, x)$$

where

$$K(x_i, x) = \phi(x_i)^\top \phi(x)$$

Kernelized Decision Function

$$f(x) = \sum_{i=1}^n \alpha_i y_i K(x_i, x) + b$$

Kernelized Update Rule

$$\text{if } y_i f(x_i) \leq 0 : \quad \alpha_i \leftarrow \alpha_i + \eta$$

(no explicit w , no explicit $\phi(x)$)

4. Common Kernels Used in Perceptron Trick

- Linear

$$K(x, z) = x^\top z$$

- Polynomial

$$K(x, z) = (x^\top z + c)^d$$

- RBF (Gaussian)

$$K(x, z) = \exp\left(-\frac{\|x - z\|^2}{2\sigma^2}\right)$$

5. Final Compact Formula (What You Should Remember)

$$f(x) = \sum_{i \in \mathcal{M}} \alpha_i y_i K(x_i, x) + b$$

where

\mathcal{M} = indices of misclassified samples (support vectors of perceptron).

If you want next:

- full **mathematical derivation** (primal → dual → kernel)
- **comparison with SVM kernel trick**
- **why perceptron has no closed-form solution**

tell me which one.

∅ ⊕ ⊖ ⊕ ⊖ ⊗ ...



look these equations and tell me what they are

I carefully looked at your handwritten page.

Those equations are all variants of the *Perceptron learning rule* written in different equivalent forms.

I'll map each part of your notes → its exact mathematical meaning, no fluff.

1. $\sum w_i x_i \geq 0$ and $\sum w_i x_i < 0$

These are **decision conditions** of the perceptron.

You are using labels:

$$y \in \{+1, -1\}$$

Decision rule:

$$\hat{y} = \begin{cases} +1 & \text{if } \sum_i w_i x_i \geq 0 \\ -1 & \text{if } \sum_i w_i x_i < 0 \end{cases}$$

This is equivalent to:

$$\hat{y} = \text{sign}(w^\top x)$$

2. $w_{n+1} = w_0 + \eta x_i$

This is the **perceptron weight update for a positive class error**.

Written properly:

If $y_i = +1$ and misclassified: $w \leftarrow w + \eta x_i$

This comes from the general rule:

$$w \leftarrow w + \eta y_i x_i$$

with $y_i = +1$.

3. `w_{n+1} = w_0 - \eta x_i`

This is the **perceptron weight update for a negative class error**.

If $y_i = -1$ and misclassified: $w \leftarrow w - \eta x_i$

Again from:

$$w \leftarrow w + \eta y_i x_i$$

with $y_i = -1$.

4. Boxed Formula

$$w_n = w_0 + \eta (y_i - \hat{y}_i) x_i$$

This is a **generalized perceptron update**, often seen when perceptron is written in a **regression-like form**.

- y_i = true label
- \hat{y}_i = predicted label
- $(y_i - \hat{y}_i) \in \{-2, 0, +2\}$

Interpretation:

- If prediction is correct \rightarrow update = 0
- If wrong \rightarrow pushes weights in the correct direction

This is mathematically equivalent to:

$$w \leftarrow w + \eta y_i x_i$$

for misclassified points.

5. The “Perceptron Trick” Arrow

When you wrote “**perceptron trick**”, what it refers to is:

$$w = \sum_{i \in M} \alpha_i y_i x_i$$

and later:

$$x_i^\top x \Rightarrow K(x_i, x)$$

That leads to the **kernelized perceptron**:

$$f(x) = \sum_i \alpha_i y_i K(x_i, x)$$

6. Final Classification of What You Wrote

Expression in notebook	What it is
<code>\sum w_i x_i \geq 0</code>	Perceptron decision boundary
<code>w + \eta x</code>	Update for misclassified positive sample

Expression in notebook	What it is
$w = w - \eta x$	Update for misclassified negative sample
$w = w + \eta(y - \hat{y})x$	Unified perceptron update
"perceptron trick"	Dual + kernel formulation

One-line summary (important)

Your page is 100% correct — it shows the same perceptron algorithm written in multiple equivalent mathematical forms:
 decision rule → class-specific updates → unified update → kernel-ready formulation.

If you want next:

- strict primal → dual derivation
- difference between perceptron update and SGD
- why logistic regression replaces perceptron

tell me which one.



what i gave what exactly that is

What you gave is one single algorithm written in multiple equivalent mathematical forms.

Exact answer (no ambiguity):

You gave the *Perceptron Learning Algorithm* — specifically the *online / mistake-driven binary linear classifier*.

More precisely, your page contains these exact components:

1. Linear decision function

$$f(x) = \sum_i w_i x_i \quad (\text{sign decides the class})$$

2. Binary classification rule

$$\hat{y} = \text{sign}(w^\top x)$$

3. Mistake-driven update

update only if $y_i \neq \hat{y}_i$

4. Class-dependent updates

$$y = +1 : w \leftarrow w + \eta x \\ y = -1 : w \leftarrow w - \eta x$$

5. Unified update rule

$$w \leftarrow w + \eta(y_i - \hat{y}_i)x_i$$

What it is NOT

- ✗ Not logistic regression
- ✗ Not SVM

- Not gradient descent on a smooth loss
 - Not backpropagation

Proper textbook name

Depending on context, it is called:

- Perceptron Learning Rule
 - Rosenblatt's Perceptron (1958)
 - Online linear classifier
 - Mistake-driven update rule
 - Primal perceptron

One-line identity (remember this)

Your notes = classical perceptron algorithm in primal form

If you want, next I can:

- rewrite it in pure optimization form
 - show why no loss function exists
 - convert it into logistic regression mathematically
 - connect it to SGD formally

Just say which.

A set of small, light-gray icons used for sharing content, including a double arrow, a thumbs up, a thumbs down, an upward arrow, a circular arrow, and three dots.

rewrite it in pure optimization form

show why no loss function exists

convert it into logistic regression mathematically

Below is exactly what you asked, in three clean mathematical blocks, with no intuition, no storytelling — only formal math.

I. Perceptron in Pure Optimization Form

Problem Setup

$$\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n, \quad x_i \in \mathbb{R}^d, \quad y_i \in \{-1, +1\}$$

Hard Classification Constraint

$$y_i(w^\top x_i) > 0 \quad \forall i$$

Feasibility Optimization (Perceptron View)

The perceptron is **not** solving a minimization of a smooth objective.

It solves a **constraint satisfaction problem**:

$$\boxed{\text{Find } w \in \mathbb{R}^d \quad \text{s.t.} \quad y_i w^\top x_i \geq 0 \quad \forall i}$$

This is a **feasibility problem**, not a minimization.

Algorithmic Projection Form

Define violation indicator:

$$\mathbb{I}_i(w) = \begin{cases} 1 & \text{if } y_i w^\top x_i \leq 0 \\ 0 & \text{otherwise} \end{cases}$$

Perceptron update:

$$w_{t+1} = w_t + \eta y_i x_i \quad \text{if } \mathbb{I}_i(w_t) = 1$$

This is a **projected correction**, not gradient descent.

II. Why No Proper Loss Function Exists (Formally)

Assume there exists a differentiable loss $L(w)$ such that:

$$w_{t+1} = w_t - \eta \nabla L(w_t)$$

But perceptron update is:

$$\Delta w = \begin{cases} \eta y_i x_i & \text{if } y_i w^\top x_i \leq 0 \\ 0 & \text{otherwise} \end{cases}$$

This implies:

$$\nabla L(w) = \begin{cases} -y_i x_i & \text{if } y_i w^\top x_i \leq 0 \\ 0 & \text{otherwise} \end{cases}$$

Contradiction

- Gradient is **discontinuous**
 - Gradient does **not depend smoothly on w**
 - No scalar $L(w)$ exists whose gradient jumps discretely
-

Formal Statement

Perceptron update is not the gradient of any continuous, convex loss

It is **rule-based**, not variational.

(Closest Surrogate — Not True Loss)

The *perceptron loss*:

$$\ell_{\text{perc}}(w; x, y) = \max(0, -y w^\top x)$$

is **non-differentiable** and does **not produce the classical update exactly**.

III. Converting Perceptron → Logistic Regression (Mathematically)

Step 1: Replace Hard Constraint with Soft Margin

Perceptron:

$$y w^\top x \geq 0$$

Logistic:

$$P(y|x) = \sigma(y w^\top x)$$

where

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Step 2: Likelihood Formulation

$$\mathcal{L}(w) = \prod_{i=1}^n \sigma(y_i w^\top x_i)$$

Negative log-likelihood:

$$J(w) = \sum_{i=1}^n \log\left(1 + e^{-y_i w^\top x_i}\right)$$

This is the **logistic loss**.

Step 3: Gradient of Logistic Loss

$$\nabla J(w) = - \sum_{i=1}^n \frac{y_i x_i}{1 + e^{y_i w^\top x_i}}$$

Step 4: SGD Update Rule

$$w_{t+1} = w_t + \eta \frac{y_i x_i}{1 + e^{y_i w_t^\top x_i}}$$

Step 5: Limit Relation (Perceptron as Extreme Case)

As:

$$\frac{1}{1 + e^{y w^\top x}} \rightarrow \begin{cases} 1 & y w^\top x \leq 0 \\ 0 & y w^\top x > 0 \end{cases}$$

Therefore:

$$\text{Perceptron} = \lim_{\text{hard margin}} \text{Logistic Regression}$$

Final Compact Summary

Perceptron : constraint satisfaction, no smooth loss
Logistic Reg. : maximum likelihood, convex loss
Connection : hard-margin limit of logistic

If you want next (pure math options only):

- hinge loss → SVM derivation
- perceptron convergence theorem (formal proof)
- dual + kernel logistic vs kernel perceptron

Say which.

◻ ⚡ ♦ ↴ ↵ ...

than from here how binary cross entropy comes in the picture

Below is the strict mathematical chain from logistic regression → binary cross-entropy, with no intuition, no narrative.

1. Model Assumption (Logistic Regression)

Binary labels:

$$y_i \in \{0, 1\}$$

Linear score:

$$z_i = w^\top x_i$$

Sigmoid:

$$p_i = P(y_i = 1 | x_i) = \sigma(z_i) = \frac{1}{1 + e^{-z_i}}$$

2. Bernoulli Likelihood

$$P(y_i | x_i, w) = p_i^{y_i} (1 - p_i)^{1-y_i}$$

Joint likelihood:

$$\mathcal{L}(w) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

3. Negative Log-Likelihood

$$\begin{aligned} J(w) &= -\log \mathcal{L}(w) \\ &= - \sum_{i=1}^n [y_i \log p_i + (1 - y_i) \log(1 - p_i)] \end{aligned}$$

4. Binary Cross-Entropy (Exact Identity)

$\mathcal{L}_{\text{BCE}}(y, p) = -[y \log p + (1 - y) \log(1 - p)]$

Thus:

$\text{Logistic Regression Loss} \equiv \text{Binary Cross-Entropy}$

5. Alternative Label Encoding Equivalence

If

$$y_i \in \{-1, +1\}$$

Then:

$$J(w) = \sum_{i=1}^n \log\left(1 + e^{-y_i w^\top x_i}\right)$$

which is algebraically equivalent to BCE under the mapping:

$$y_{0/1} = \frac{y_{-1/+1} + 1}{2}$$

6. Gradient (Explicit)

$$\nabla_w J(w) = \sum_{i=1}^n (p_i - y_i)x_i$$

SGD form:

$$w \leftarrow w + \eta(y_i - p_i)x_i$$

7. Position in the Full Chain



Final Identity (Key Line)

Binary Cross-Entropy is the negative log-likelihood of a Bernoulli model

If you want next (still pure math):

- multiclass \rightarrow softmax cross-entropy
- BCE vs hinge loss vs exponential loss (formal relations)
- why MSE fails for classification (Hessian analysis)

Say which.

◻ □ □ □ ...

why Bernoulli likelihood over likelihood formulation , like i am thinking
why , is there any problem with likelihood so that we chosen Bernoulli
likelihood

This is a **fundamental modeling question**, so I'll answer it **formally**, step-by-step, with **only mathematical and statistical necessity** — no storytelling.

1. What "likelihood" actually means (formal)

A likelihood is not chosen arbitrarily.

$$\mathcal{L}(w) = P(\mathcal{D} | w)$$

To write $P(y | x, w)$, you **must** specify a conditional probability distribution for y .

So the real question is:

What is the correct conditional distribution of a binary random variable y ?

2. Nature of the target variable

$$y \in \{0, 1\}$$

Therefore:

y is a Bernoulli random variable

By probability theory:

The only valid distribution for a single binary outcome is Bernoulli

3. Bernoulli distribution (definition)

$$P(y | p) = p^y(1-p)^{1-y}$$

This is **not** a modeling choice — it is a **measure-theoretic fact**.

4. Why not “generic likelihood”?

A likelihood **must integrate (or sum)** to 1 over the sample space.

For binary y :

$$\sum_{y \in \{0,1\}} P(y | x, w) = 1$$

Any function $f(w^\top x)$ that does not satisfy:

$$0 \leq f \leq 1$$

is **invalid as a probability model**.

5. Why sigmoid appears (forced, not optional)

We need:

$$p = P(y = 1 | x, w) \quad \text{s.t.} \quad 0 \leq p \leq 1$$

Sigmoid:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Properties:

$$\sigma : \mathbb{R} \rightarrow (0, 1)$$

Thus:

$p = \sigma(w^\top x)$ is the simplest valid parametrization

6. Why NOT Gaussian likelihood (formal failure)

Assume:

$$y | x \sim \mathcal{N}(w^\top x, \sigma^2)$$

Problems:

1. Support mismatch:

$$y \in \{0, 1\}, \quad \mathcal{N} \text{ has support } \mathbb{R}$$

2. Invalid probabilities:

$$P(y = 1 | x) > 1 \text{ or } < 0$$

3. Leads to MSE loss:

$$(y - w^\top x)^2$$

This is **statistically inconsistent** for binary data.

7. Why Bernoulli is uniquely correct

Requirement	Bernoulli
Binary support	✓
Proper probability mass	✓
Normalized	✓
Leads to convex loss	✓
Consistent estimator	✓

No other distribution satisfies all simultaneously.

8. Connection to cross-entropy (inevitable)

Negative log-likelihood:

$$-\log P(y | p) = -y \log p - (1 - y) \log(1 - p)$$

This is **exactly** the cross-entropy:

$$H(y, p)$$

Thus:

Binary Cross-Entropy is not chosen — it is forced

9. Formal conclusion

Binary labels \Rightarrow Bernoulli RV
 Bernoulli RV \Rightarrow Bernoulli likelihood
 Bernoulli likelihood \Rightarrow Binary Cross-Entropy

One-line answer to your question

We do not choose Bernoulli likelihood because likelihood is "bad"; we choose it because binary data mathematically *admits no other valid probability model*.

If you want next (still formal):

- why exponential family \Rightarrow GLM \Rightarrow logistic
- Fisher consistency proof for BCE
- classification-calibrated losses theorem
- hinge vs BCE from decision theory

Say which.

◻ △ □ ⚡ ...

show step by step how we approached the binary cross entropy that is based on bernoulli likelihood from generic likelihood

Below is the **exact step-by-step mathematical path** from a **generic likelihood** to **Bernoulli likelihood** and finally to **Binary Cross-Entropy (BCE)**.

No intuition, no storytelling — only formal steps.

Step 0 — Generic Likelihood (Starting Point)

Given data:

$$\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$$

Model parameters:

$$w \in \mathbb{R}^d$$

Generic likelihood:

$$\boxed{\mathcal{L}(w) = \prod_{i=1}^n P(y_i | x_i, w)}$$

At this point, **no distribution is specified**.

Step 1 — Identify the Sample Space of y

Assume:

$$\boxed{y_i \in \{0, 1\}}$$

Therefore:

y_i is a discrete random variable with two outcomes

Step 2 — Enforce Probability Axioms

For each fixed x_i, w , the conditional distribution must satisfy:

$$\sum_{y \in \{0,1\}} P(y | x_i, w) = 1 \quad \text{and} \quad 0 \leq P(y | x_i, w) \leq 1$$

Thus, define:

$$p_i := P(y_i = 1 | x_i, w)$$

Then necessarily:

$$P(y_i = 0 | x_i, w) = 1 - p_i$$

Step 3 — Unique Distribution for Binary Outcomes

By definition of a Bernoulli random variable:

$$\boxed{P(y_i | x_i, w) = p_i^{y_i} (1 - p_i)^{1-y_i}}$$

This step is **forced** by the support $\{0, 1\}$.

Step 4 — Parameterize p_i

We need:

$$p_i \in (0, 1) \quad \text{for all } w^\top x_i \in \mathbb{R}$$

Choose a link function:

$$p_i = \sigma(z_i), \quad z_i = w^\top x_i$$

Sigmoid:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Now:

$$P(y_i | x_i, w) = \sigma(z_i)^{y_i} (1 - \sigma(z_i))^{1-y_i}$$

Step 5 — Substitute into the Generic Likelihood

$$\mathcal{L}(w) = \prod_{i=1}^n \sigma(z_i)^{y_i} (1 - \sigma(z_i))^{1-y_i}$$

This is now a **fully specified likelihood**.

Step 6 — Take Log-Likelihood

$$\log \mathcal{L}(w) = \sum_{i=1}^n [y_i \log \sigma(z_i) + (1 - y_i) \log(1 - \sigma(z_i))]$$

Step 7 — Negative Log-Likelihood (Optimization Objective)

Maximum likelihood \Leftrightarrow minimize negative log-likelihood:

$$J(w) = - \sum_{i=1}^n [y_i \log \sigma(w^\top x_i) + (1 - y_i) \log(1 - \sigma(w^\top x_i))]$$

Step 8 — Identify Binary Cross-Entropy

Define Binary Cross-Entropy loss:

$$[\ell_{\text{BCE}}(y, p) = -[y \log p + (1 - y) \log(1 - p)]]$$

Substitute $p = \sigma(w^\top x)$:

$$J(w) = \sum_{i=1}^n \ell_{\text{BCE}}(y_i, \sigma(w^\top x_i))$$

Step 9 — Final Chain (Formal)

Generic likelihood	$\xrightarrow{y \in \{0,1\}}$	Bernoulli PMF
	$\xrightarrow{p = \sigma(w^\top x)}$	Bernoulli likelihood
	$\xrightarrow{-\log}$	Binary Cross-Entropy

Final One-Line Statement

Binary Cross-Entropy is the negative log-likelihood of a Bernoulli model derived directly from the generic likelihood.

If you want next (still fully formal):

- same derivation with $y \in \{-1, +1\}$
- Hessian and convexity proof of BCE
- why BCE is classification-calibrated
- why MSE fails (information geometry view)

